

Statistical modelling with Python

Simple linear Regression with Python-native SKiKit-Learn Package

Imports necessary for Regression models with SciKitLearn

```
In [2]:  
  
# imports necessary for statistical modelling  
# convention is to define imports at the top of scripts  
  
import numpy as np  
import pandas as pd  
import seaborn as sns  
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score
```

Using the ALLBUS data set for standard linear regressions

```
In [3]:  
  
# import data set, set directory  
home_dir = "/home/tobias_giesemann/Dropbox/Uni_Master/02SS19/01Advanced_Statistical_Modelling/Essay/"  
  
# using pandas dataframe as a similar data structure to R dataframes  
allbus_df = pd.read_csv(home_dir+"data/allbus_small.csv", index_col=0)  
  
#show head  
allbus_df.head()
```

Out[3]:

	income	sex	age	edueyears	eastwest	socialclass_self	leftright
1	1800.0	FRAU	47.0	13.0	NEUE BUNDESALAENDER	MITTELSCHICHT	2.0
2	2000.0	MANN	52.0	13.0	NEUE BUNDESALAENDER	MITTELSCHICHT	4.0
3	2500.0	MANN	61.0	9.0	ALTE BUNDESALAENDER	MITTELSCHICHT	8.0
4	860.0	FRAU	54.0	12.0	ALTE BUNDESALAENDER	NaN	3.0
5	NaN	MANN	71.0	NaN	ALTE BUNDESALAENDER	OBERSCHICHT	7.0

```
In [4]:  
  
#description of numeric variables  
allbus_df.describe()
```


Out[4]:

	income	age	edueyears	leftright
count	2654.000000	3486.000000	3303.000000	3335.000000
mean	1609.666164	51.143144	12.632758	5.078561
std	1100.712568	17.567575	3.694805	1.700014
min	1.000000	18.000000	4.000000	1.000000
25%	850.000000	37.000000	10.000000	4.000000
50%	1400.000000	52.000000	12.000000	5.000000
75%	2000.000000	65.000000	15.000000	6.000000
max	9500.000000	97.000000	33.000000	10.000000

Plotting in Python with package "Seaborn"


In [13]:

```
# basic plotting with seaborn plotting package
basic_plot = sns.relplot(x="eduyears",y="income", data=allbus_df)
```




In [12]:

```
# little more advanced plotting
fair_plot = sns.relplot(x="eduyears",y="income", hue="sex", data=allbus_df)
```




In [11]:

```
# too much for one plot, but great features
overly_advanced_plot = sns.relplot(x="eduyears",
    y="income",
    hue="leftright",
    style="sex",
    size="age",
    palette="ch:r=-.5,l=.75",
    sizes=(10,300),
    data=allbus_df)
overly_advanced_plot.savefig(home_dir+'figures/python/overly_advanced_plot.pdf')
```



In [10]:

```
# simple and fast plot for linear models -> very powerful
lm_plot = sns.lmplot('edueyears', 'income', data=allbus_df, fit_reg=True)
```



Linear Models with SciKit Learn

In [4]:

```
# reshape input variables
regression_df1 = allbus_df[['income', 'age']]
regression_df1.dropna(inplace=True)
X = np.array([regression_df1.age]).reshape(-1, 1)
y = np.array([regression_df1.income]).reshape(-1, 1)

# define model
regression_model = LinearRegression()
regression_model.fit(X, y)

y_predicted = regression_model.predict(X)

# model evaluation
mse = mean_squared_error(y, y_predicted)
r2 = r2_score(y, y_predicted)

# printing values
print('Estimate Std.: ', regression_model.coef_)
print('Intercept:', regression_model.intercept_)
print('Mean squared error: ', mse)
print('R2 score: ', r2)
```

```
Estimate Std.: [[2.79715344]]
Intercept: [1465.57232066]
Mean squared error: 1208951.0226202777
R2 score: 0.002015611062190281
```

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

Linear Regression with Package StatsModels

In [23]:

```
import statsmodels.api as sm

# import data set, set directory
home_dir = "/home/tobias_gieseemann/Dropbox/Uni_Master/02SS19/01Advanced_Statistical_Modelling/Essay/"

# using pandas dataframe as a similar data structure to R dataframes
allbus_df = pd.read_csv(home_dir+"data/allbus_small.csv", index_col=0)

# show head
allbus_df.head()

# drop nan values
allbus_df.dropna(inplace = True)
```

Linear Regression

As we can see, this emulation of the R syntax is much easier to implement

In [25]:

```
import statsmodels.formula.api as smf

lm1 = smf.ols('income ~ 1+age+sex+eduyears+eastwest+leftright', data=allbus_df).fit()
print(lm1.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          income      R-squared:                0.218
Model:                  OLS        Adj. R-squared:             0.216
Method:                 Least Squares      F-statistic:           134.5
Date:                  Thu, 25 Jul 2019    Prob (F-statistic):     4.95e-126
Time:                  16:40:57          Log-Likelihood:        -20109.
No. Observations:      2418             AIC:                  4.023e+04
Df Residuals:          2412             BIC:                  4.026e+04
Df Model:              5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-412.2895	128.587	-3.206	0.001	-664.442	-160.137
sex[T.MANN]	647.4251	40.427	16.015	0.000	568.149	726.701
eastwest[T.NEUE BUNDESLAENDER]	-371.3275	42.863	-8.663	0.000	-455.380	-287.275
age	8.8806	1.188	7.477	0.000	6.552	11.210
eduyears	99.1138	5.674	17.468	0.000	87.987	110.240
leftright	27.9986	11.870	2.359	0.018	4.722	51.275

```
=====
Omnibus:                730.914      Durbin-Watson:           1.999
Prob(Omnibus):          0.000        Jarque-Bera (JB):        3151.537
Skew:                   1.402        Prob(JB):                0.00
Kurtosis:               7.839        Cond. No.                356.
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Residual Plot with Python

As this plot is often needed to check for homoscedascity, I will check if there is an easy implementation in python as well. As we will see, this is not the case, and we rather need to build our own plot.


In [41]:

```
# fitted values (need a constant term for intercept)
model_fitted_y = lm1.fittedvalues

# residual plot
sns.residplot(x=model_fitted_y,
              y="income",
              data=allbus_df,
              lowess = True,
              scatter_kws={'alpha':0.5},
              line_kws={"color": "red", "lw": 1, "alpha": 0.9})
```

Out[41]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f0e6a0f6438>



Here, we would rather have to check for heteroscedascity and maybe consider a quantile-regression model