

Generalized Linear Models

Logit Models

In [61]:

```
import pandas as pd
import seaborn as sns
from statsmodels.formula.api import logit, glm
import statsmodels.api as sm

from plotnine import *

# import data set, set directory
home_dir = "/home/tobias_gieseemann/Dropbox/Uni_Master/02SS19/01Advanced_Statistical_Modelling/Essay/"

# using pandas dataframe as a similar data structure to R dataframes
allbus_df = pd.read_csv(home_dir+"data/allbus_full.csv")
print("full data set: ", allbus_df.shape)
allbus_df = allbus_df[["eastwest", "sex", "age", "lt15"]]
allbus_df.contacts = allbus_df.lt15
allbus_df["any_contact"] = 0
allbus_df.loc[allbus_df.contacts > 0, "any_contact"] = 1

print(allbus_df.contacts.unique())
print(allbus_df.any_contact.unique())

#show head
allbus_df.head()
```

full data set: (3490, 793)
[nan 1. 2. 4. 3. 6. 5. 9. 8. 12. 7. 18. 11. 10.]
[0 1]

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/ipykernel_launcher.py:16: UserWarning: Pandas doesn't allow column s to be created via a new attribute name - see <https://pandas.pydata.org/pandas-docs/stable/indexing.html#attribute-access>
app.launch_new_instance()

Out[61]:

	eastwest	sex	age	lt15	any_contact
0	NEUE BUNDESLAENDER	FRAU	47.0	NaN	0
1	NEUE BUNDESLAENDER	MANN	52.0	NaN	0
2	ALTE BUNDESLAENDER	MANN	61.0	1.0	1
3	ALTE BUNDESLAENDER	FRAU	54.0	NaN	0
4	ALTE BUNDESLAENDER	MANN	71.0	NaN	0

In [58]:

```
allbus_df["age_cat"] = pd.cut(allbus_df.age, [17,30,40,50,60,70,100])
```

```
(ggplot(data = allbus_df,
        mapping = aes(x = "age_cat",
                      y = "lt15"))
+geom_point(mapping=aes(color=allbus_df["lt15"])))
)
```

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/mizani/bounds.py:352: RuntimeWarning: invalid value encountered in less

```
    outside = (x < range[0]) | (x > range[1])
```

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/mizani/bounds.py:352: RuntimeWarning: invalid value encountered in greater

```
    outside = (x < range[0]) | (x > range[1])
```

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/matplotlib/colors.py:527: RuntimeWarning: invalid value encountered in less

```
    xa[xa < 0] = -1
```


/home/tobias_gieseemann/.local/lib/python3.6/site-packages/plotnine/layer.py:449: UserWarning: geom_point : Removed 1756 rows containing missing values.

```
    self.data = self.geom.handle_na(self.data)
```

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/plotnine/layer.py:553: MatplotlibDeprecationWarning:

The iterable function was deprecated in Matplotlib 3.1 and will be removed in 3.3. Use np.iterable instead.

```
    return cbook.iterable(var) and not is_string(var)
```



Out[58]:

```
<ggplot: (8771815829859)>
```

Probit Modell

In [79]:

```
# Probit Modell als GLM
```

```
probit_model = glm("any_contact ~ age + sex + eastwest", data = allbus_df, family = sm.families.Binomial(link=sm.families.links.probit)).fit()
print(probit_model.summary())
```

Generalized Linear Model Regression Results

```
=====
Dep. Variable:          any_contact      No. Observations:          3486
Model:                  GLM              Df Residuals:              3482
Model Family:           Binomial         Df Model:                  3
Link Function:           probit           Scale:                    1.0000
Method:                  IRLS             Log-Likelihood:             -2288.8
Date:                   Thu, 25 Jul 2019   Deviance:                   4577.6
Time:                   22:18:47           Pearson chi2:               3.49e+03
No. Iterations:         5
Covariance Type:        nonrobust
=====
```

	coef	std err	z	P> z	[0.025	0.975]
Intercept	0.7664	0.071	10.745	0.000	0.627	0.906
sex[T.MANN]	-0.0074	0.043	-0.170	0.865	-0.092	0.078
eastwest[T.NEUE BUNDESLAENDER]	-0.5642	0.046	-12.145	0.000	-0.655	-0.473
age	-0.0114	0.001	-9.150	0.000	-0.014	-0.009

/home/tobias_gieseemann/.local/lib/python3.6/site-packages/ipykernel_launcher.py:4: DeprecationWarning: Calling Family(..) with a link class as argument is deprecated.

Use an instance of a link class instead.

after removing the cwd from sys.path.

Logit Modell

In [88]:

```
#selbiges Modell wie oben
logit_model = glm("any_contact ~ age + sex + eastwest", data = allbus_df,
                  family = sm.families.Binomial(link=sm.families.links.logit # or None, da logit default für Binomialfamilie
                  )).fit()

print(logit_model.summary())

print("Maximal beobachtbare Effekte der Einflussgrößen: \n",0.25*logit_model.params)
print("Multiplikative Effektinterpretation: \n", (np.exp(logit_model.params)-1)*100)

# Margins konnten leider nicht herausgearbeitet werden, da ist die Dokumentation unvollständig,
# allerdings funktioniert es für das logit modell über einen Umweg:

logit_model2 = logit("any_contact ~ age + sex + eastwest", allbus_df).fit()
print(logit_model2.summary())

margins = logit_model2.get_margeff().summary_frame()
print(margins)
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	any_contact	No. Observations:	3486			
Model:	GLM	Df Residuals:	3482			
Model Family:	Binomial	Df Model:	3			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-2288.7			
Date:	Thu, 25 Jul 2019	Deviance:	4577.5			
Time:	22:33:19	Pearson chi2:	3.49e+03			
No. Iterations:	4					
Covariance Type:	nonrobust					
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	1.2419	0.117	10.651	0.000	1.013	1.470
sex[T.MANN]	-0.0115	0.070	-0.163	0.871	-0.149	0.126
eastwest[T.NEUE BUNDESLAENDER]	-0.9089	0.076	-12.021	0.000	-1.057	-0.761
age	-0.0186	0.002	-9.120	0.000	-0.023	-0.015
=====						

Maximal beobachtbare Effekte der Einflussgrößen:

Intercept	0.310475
sex[T.MANN]	-0.002864
eastwest[T.NEUE BUNDESLAENDER]	-0.227232
age	-0.004643

dtype: float64

Multiplikative Effektinterpretation:

Intercept	246.218916
sex[T.MANN]	-1.139177
eastwest[T.NEUE BUNDESLAENDER]	-59.704452
age	-1.840169

dtype: float64

Optimization terminated successfully.
Current function value: 0.656554
Iterations 5

Logit Regression Results						
=====						
Dep. Variable:	any_contact	No. Observations:	3486			
Model:	Logit	Df Residuals:	3482			
Method:	MLE	Df Model:	3			
Date:	Thu, 25 Jul 2019	Pseudo R-squ.:	0.05277			
Time:	22:33:20	Log-Likelihood:	-2288.7			
converged:	True	LL-Null:	-2416.3			
Covariance Type:	nonrobust	LLR p-value:	5.390e-55			
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	1.2419	0.117	10.651	0.000	1.013	1.470
sex[T.MANN]	-0.0115	0.070	-0.163	0.871	-0.149	0.126
eastwest[T.NEUE BUNDESLAENDER]	-0.9089	0.076	-12.021	0.000	-1.057	-0.761
age	-0.0186	0.002	-9.120	0.000	-0.023	-0.015
=====						

	dy/dx	Std. Err.	z	Pr(> z)	\
sex[T.MANN]	-0.002660	0.016330	-0.162902	8.705953e-01	
eastwest[T.NEUE BUNDESLAENDER]	-0.211037	0.016105	-13.104195	3.115508e-39	
age	-0.004312	0.000451	-9.559622	1.181875e-21	

	Conf. Int. Low	Cont. Int. Hi.
sex[T.MANN]	-0.034666	0.029345
eastwest[T.NEUE BUNDESLAENDER]	-0.242602	-0.179473
age	-0.005196	-0.003428

/home/tobias_giesemann/.local/lib/python3.6/site-packages/ipykernel_launcher.py:1: DeprecationWarning: Calling Family(..) with a link class as argument is deprecated.
Use an instance of a link class instead.
"""Entry point for launching an IPython kernel.

Poisson -Regression

In [99]:

```
#selbiges Modell wie oben
logit_model = glm("lt15 ~ age + sex + eastwest", data = allbus_df,
                  family = sm.families.Poisson()).fit()
print(logit_model.summary())

print("Multiplikative Effektinterpretation: \n", (np.exp(logit_model.params)-1)*100)
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	lt15	No. Observations:	1733			
Model:	GLM	Df Residuals:	1729			
Model Family:	Poisson	Df Model:	3			
Link Function:	log	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-2858.5			
Date:	Thu, 25 Jul 2019	Deviance:	1447.2			
Time:	23:28:51	Pearson chi2:	1.85e+03			
No. Iterations:	4					
Covariance Type:	nonrobust					
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	0.9280	0.052	17.816	0.000	0.826	1.030
sex[T.MANN]	0.0219	0.034	0.650	0.516	-0.044	0.088
eastwest[T.NEUE BUNDESLAENDER]	-0.1311	0.042	-3.150	0.002	-0.213	-0.050
age	-0.0043	0.001	-4.381	0.000	-0.006	-0.002
=====						
Multiplikative Effektinterpretation:						
Intercept	152.942160					
sex[T.MANN]	2.218618					
eastwest[T.NEUE BUNDESLAENDER]	-12.290349					
age	-0.428598					
dtype:	float64					

Hier funktioniert der Trick nicht, ein einfacheres Modell zur Margin Schätzung zu verwenden. Die müsste man dann wohl zu Fuß berechnen.

Negativ-Binomial- Regression

In [101]:

```
#selbiges Modell wie oben
logit_model = glm("lt15 ~ age + sex + eastwest", data = allbus_df,
                  family = sm.families.NegativeBinomial()).fit()
print(logit_model.summary())

print("Multiplikative Effektinterpretation: \n", (np.exp(logit_model.params)-1)*100)
```

Generalized Linear Model Regression Results						
=====						
Dep. Variable:	lt15	No. Observations:	1733			
Model:	GLM	Df Residuals:	1729			
Model Family:	NegativeBinomial	Df Model:	3			
Link Function:	log	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-3323.1			
Date:	Thu, 25 Jul 2019	Deviance:	438.23			
Time:	23:29:41	Pearson chi2:	610.			
No. Iterations:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	z	P> z	[0.025	0.975]

Intercept	0.9326	0.092	10.179	0.000	0.753	1.112
sex[T.MANN]	0.0227	0.059	0.386	0.699	-0.093	0.138
eastwest[T.NEUE BUNDESLAENDER]	-0.1338	0.071	-1.892	0.059	-0.272	0.005
age	-0.0044	0.002	-2.584	0.010	-0.008	-0.001
=====						
Multiplikative Effektinterpretation:						
Intercept	154.116026					
sex[T.MANN]	2.297112					
eastwest[T.NEUE BUNDESLAENDER]	-12.522152					
age	-0.437916					
dtype:	float64					